

SQL Lecture II

G51DBI – Databases and Interfaces

Yorgos Tzimiropoulos

yorgos.tzimiropoulos@nottingham.ac.uk

Overview of weeks 2-4

We will see how to translate English to Relational Algebra and SQL queries **and** vice versa

English: “Find all universities with > 20000 students”

Relational Algebra: $\pi_{uName}(\sigma_{enr > 20000}(University))$

SQL: **Select** uName **From** University **Where** University.enr>20000

Theory is easy and simple

But a sequence of simple operations is not always so obvious!

2

This Lecture

➤ SQL SELECT

- WHERE Clauses
- SELECT from multiple tables

➤ More SQL SELECT

- Aliases
- ‘Self-Joins’
- Subqueries
- IN, EXISTS, ANY, ALL

3

SQL SELECT Overview

SELECT

[**DISTINCT** | **ALL**] column-list

FROM table-names

[**WHERE** condition]

[**ORDER BY** column-list]

[**GROUP BY** column-list]

[**HAVING** condition]

([] optional, | or)

4

Example Tables

Student

ID	First	Last
S103	John	Smith
S104	Mary	Jones
S105	Jane	Brown
S106	Mark	Jones
S107	John	Brown

Course

Code	Title
DBS	Database Systems
PR1	Programming 1
PR2	Programming 2
IAI	Introduction to AI

Grade

ID	Code	Mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

5

DISTINCT and ALL

- Sometimes you end up with duplicate entries
- Using **DISTINCT** removes duplicates
- Using **ALL** retains duplicates
- **ALL** is used as a default if neither is supplied
- These will work over multiple columns

SELECT ALL Last
FROM Student;

Last
Smith
Jones
Brown
Jones
Brown

SELECT DISTINCT Last
FROM Student;

Last
Smith
Jones
Brown

6

WHERE Clauses

- In most cases returning all the rows is not necessary
 - A WHERE clause restricts rows that are returned
 - It takes the form of a condition – only rows that satisfy the condition are returned
- Example conditions:
 - Mark < 40
 - First = 'John'
 - First <> 'John'
 - First = Last
 - (First = 'John') AND (Last = 'Smith')
 - (Mark < 40) OR (Mark > 70)

7

WHERE Examples

```
SELECT * FROM Grade
WHERE Mark >= 60;

SELECT DISTINCT ID
FROM Grade
WHERE Mark >= 60;
```

8

WHERE Examples

```
SELECT * FROM Grade
WHERE Mark >= 60;

SELECT DISTINCT ID
FROM Grade
WHERE Mark >= 60;
```

ID	Code	Mark
S103	DBS	72
S104	PR1	68
S104	IAI	65
S107	PR1	76
S107	PR2	60

ID
S103
S104
S107

9

WHERE Examples

- Given the table:
- Write an SQL query to find a list of the ID numbers and Marks for students who have passed (scored 50% or more) in IAI

Grade		
ID	Code	Mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

ID	Mark
S103	58
S104	65

10

Solution

```
SELECT ID, Mark FROM Grade
WHERE (Code = 'IAI')
AND (Mark >= 50);
```

11

WHERE Examples

- Given the table:
- Write an SQL query to find a list of the ID numbers and Marks for students who have passed with Merit (Marks in [60, 69])

Grade		
ID	Code	Mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

ID	Mark
S104	68
S104	65
S107	60

12

Solution

```
SELECT ID, Mark FROM Grade
WHERE (Mark >=60
AND Mark < 70);
```

13

Solution (only in MySQL!)

```
SELECT ID, Mark FROM Grade WHERE
Mark BETWEEN 60 AND 69;
```

14

WHERE Examples

- Given the table:

ID	Code	Mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

- Write an SQL query to find a list of the students IDs for both the IAI and PR2 modules

ID
S103
S104
S106
S107
S107

15

Solution

```
SELECT ID FROM Grade
WHERE (Code = 'IAI'
OR Code = 'PR2');
```

16

SELECT from Multiple Tables

- Often you need to combine information from two or more tables
- You can produce the effect of a Cartesian product using:
`SELECT * FROM Table1, TableName.ColumnName Table2`
- If the tables have columns with the same name, ambiguity will result
- This can be resolved by referencing columns with the table name:

17

SELECT from Multiple Tables

```
SELECT
    First, Last, Mark
FROM
    Student, Grade
WHERE
    (Student.ID = Grade.ID)
    AND (Mark >= 40);
```

ID	First	Last
S103	John	Smith
S104	Mar	
S105	Jane	
S106	Mar	
S107	John	

ID	Code	Mark
S103	DBS	72
S103	IAI	58
S104	PR1	68
S104	IAI	65
S106	PR2	43
S107	PR1	76
S107	PR2	60
S107	IAI	35

18

SELECT from Multiple Tables

SELECT ... FROM Student, Grade WHERE ...

ID	First	Last	ID	Code	Mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S103	John	Smith	S104	PR1	68
S103	John	Smith	S104	IAI	65
S103	John	Smith	S106	PR2	43
S103	John	Smith	S107	PR1	76
S103	John	Smith	S107	PR2	60
S103	John	Smith	S107	IAI	35
S104	Mary	Jones	S103	DBS	72
S104	Mary	Jones	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65

19

SELECT from Multiple Tables

SELECT ... FROM Student, Grade
WHERE (Student.ID = Grade.ID) AND ...

ID	First	Last	ID	Code	Mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S106	Mark	Jones	S106	PR2	43
S107	John	Brown	S107	PR1	76
S107	John	Brown	S107	PR2	60
S107	John	Brown	S107	IAI	35

20

SELECT from Multiple Tables

SELECT ... FROM Student, Grade
WHERE (Student.ID = Grade.ID) AND (Mark >= 40)

ID	First	Last	ID	Code	Mark
S103	John	Smith	S103	DBS	72
S103	John	Smith	S103	IAI	58
S104	Mary	Jones	S104	PR1	68
S104	Mary	Jones	S104	IAI	65
S106	Mark	Jones	S106	PR2	43
S107	John	Brown	S107	PR1	76
S107	John	Brown	S107	PR2	60

21

SELECT from Multiple Tables

SELECT First, Last, Mark FROM Student, Grade
WHERE (Student.ID = Grade.ID) AND (Mark >= 40)

First	Last	Mark
John	Smith	72
John	Smith	58
Mary	Jones	68
Mary	Jones	65
Mark	Jones	43
John	Brown	76
John	Brown	60

22

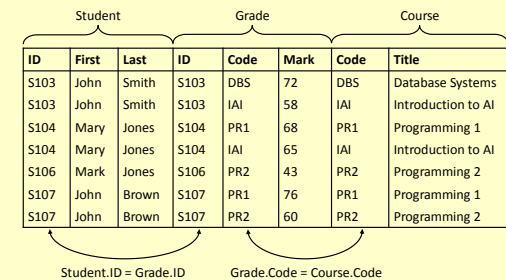
SELECT from Multiple Tables

- When selecting from multiple tables, it is almost always best to use a **WHERE** clause to find common values

```
SELECT *
FROM
  Student, Grade,
  Course
WHERE
  Student.ID =
  Grade.ID
  AND
  Course.Code =
  Grade.Code
```

23

SELECT from Multiple Tables



24

Examples

Student

sID	sName	sAddress	sYear
1	Smith	5 Arnold Close	2
2	Brooks	7 Holly Avenue	2
3	Anderson	15 Main Street	3
4	Evans	Flat 1a, High Street	2
5	Harrison	Newark Hall	1
6	Jones	Southwell Hall	1

Module

mCode	mCredits	mTitle
G51DBS	10	Database Systems
G51PRG	20	Programming
G51IAI	10	Artificial Intelligence
G52ADS	10	Algorithms

Enrolment

sID	mCode
1	G52ADS
2	G52ADS
5	G51DBS
5	G51PRG
5	G51IAI
4	G52ADS
6	G51PRG
6	G51IAI

25

Examples

- Write SQL statements to do the following:
 - Produce a list of all student names and all their enrolments (module codes)
 - Find a list of module titles being taken by the student named "Harrison"
 - Find a list of module codes and titles for all modules currently being taken by first year students

26

Solutions

- `SELECT sName, mCode FROM Student, Enrolment WHERE Student.sID = Enrolment.sID;`
- `SELECT mTitle FROM Module, Student, Enrolment WHERE (Module.mCode = Enrolment.mCode) AND (Student.sID = Enrolment.sID) AND Student.sName = "Harrison";`
- `SELECT Module.mCode, mTitle FROM Enrolment, Module, Student WHERE (Module.mCode = Enrolment.mCode) AND (Student.sID = Enrolment.sID) AND sYear = 1;`

27

Writing Queries

- When writing queries
 - There are often many ways to accomplish the same query
 - Be concerned with correctness, clarity and conciseness, in that order
 - Do not worry hugely about being clever or efficient
- Most DBMSs have query optimisers
 - Will optimise your query to improve efficiency
 - Simpler queries are easier to optimise

28

SQL SELECT Overview

```

SELECT
  [DISTINCT | ALL] column-list
FROM table-names
[WHERE condition]
[ORDER BY column-list]
[GROUP BY column-list]
[HAVING condition]
                                ([ ] optional, | or)
    
```

29

Aliases

- Aliases rename columns or tables
 - Can make names more meaningful
 - Can shorten names, making them easier to use
 - Can resolve ambiguous names
 - Two forms:
 - Column alias
`SELECT column [AS] new-col-name`
 - Table alias
`SELECT * FROM table [AS] new-table-name`
- ([] optional)

30

Alias Example

Employee

ID	Name
123	John
124	Mary

WorksIn

ID	Department
123	Marketing
124	Sales
124	Marketing

```
SELECT
  E.ID AS empID,
  E.Name, W.Department
FROM
  Employee E,
  WorksIn W
WHERE
  E.ID = W.ID;
```

Note: You cannot use a column alias in a WHERE clause

31

Alias Example

empID	Name	Department
123	John	Marketing
124	Mary	Sales
124	Mary	Marketing

```
SELECT
  E.ID AS empID,
  E.Name, W.Department
FROM
  Employee E,
  WorksIn W
WHERE
  E.ID = W.ID;
```

32

Aliases and 'Self-Joins'

- Aliases can be used to copy a table, so that it can be combined with itself:

```
SELECT A.Name FROM
  Employee A,
  Employee B
WHERE A.Dept = B.Dept
AND B.Name = 'Andy';
```

Employee

Name	Dept
John	Marketing
Mary	Sales
Peter	Sales
Andy	Marketing
Anne	Marketing

- Find the names of all employees who work in the same department as Andy.

33

Aliases and 'Self-Joins'

Employee A

A	Name	Dept
	John	Marketing
	Mary	Sales
	Peter	Sales
	Andy	Marketing
	Anne	Marketing

Employee B

B	Name	Dept
	John	Marketing
	Mary	Sales
	Peter	Sales
	Andy	Marketing
	Anne	Marketing

34

Aliases and 'Self-Joins'

```
SELECT ... FROM Employee A, Employee B ...
```

A.Name	A.Dept	B.Name	B.Dept
John	Marketing	John	Marketing
Mary	Sales	John	Marketing
Peter	Sales	John	Marketing
Andy	Marketing	John	Marketing
Anne	Marketing	John	Marketing
John	Marketing	Mary	Sales
Mary	Sales	Mary	Sales
Peter	Sales	Mary	Sales
Andy	Marketing	Mary	Sales
Anne	Marketing	Mary	Sales

35

Aliases and 'Self-Joins'

```
SELECT ... FROM Employee A, Employee B
WHERE A.Dept = B.Dept ...
```

A.Name	A.Dept	B.Name	B.Dept
John	Marketing	John	Marketing
Andy	Marketing	John	Marketing
Anne	Marketing	John	Marketing
Mary	Sales	Mary	Sales
Peter	Sales	Mary	Sales
Mary	Sales	Peter	Sales
Peter	Sales	Peter	Sales
John	Marketing	Andy	Marketing
Andy	Marketing	Andy	Marketing
Anne	Marketing	Andy	Marketing

36

Aliases and 'Self-Joins'

```
SELECT ... FROM Employee A, Employee B
WHERE A.Dept = B.Dept AND B.Name = 'Andy';
```

A.Name	A.Dept	B.Name	B.Dept
John	Marketing	Andy	Marketing
Andy	Marketing	Andy	Marketing
Anne	Marketing	Andy	Marketing

37

Aliases and 'Self-Joins'

```
SELECT A.Name FROM Employee A, Employee B
WHERE A.Dept = B.Dept AND B.Name = 'Andy';
```

A.Name
John
Andy
Anne

- Names of all employees who work in the same department as Andy.

38

Subqueries

- A SELECT statement can be nested inside another query to form a subquery
- The results of the subquery are passed back to the containing query
- For example, retrieve a list of names of people who are in Andy's department:

```
SELECT Name
FROM Employee
WHERE Dept =
(SELECT Dept
FROM Employee
WHERE Name = 'Andy');
```

39

Subqueries

```
SELECT Name
FROM Employee
WHERE Dept =
(SELECT Dept
FROM Employee
WHERE
Name = 'Andy');
```

- First the subquery is evaluated, returning 'Marketing'
- This value is passed to the main query

```
SELECT Name
FROM Employee
WHERE Dept =
'Marketing';
```

40

Subqueries

- Often a subquery will return a set of values rather than a single value
- We cannot directly compare a single value to a set. Doing so will result in an error
- Options for handling sets
 - IN – checks to see if a value is in a set
 - EXISTS – checks to see if a set is empty
 - ALL/ANY – checks to see if a relationship holds for every/one member of a set
 - NOT can be used with any of the above

41

IN

- Using IN we can see if a given value is in a set of values
- NOT IN checks to see if a given value is not in the set
- The set can be given explicitly or can be produced in a subquery

```
SELECT columns
FROM tables
WHERE value
IN set;

SELECT columns
FROM tables
WHERE value
NOT IN set;
```

42

IN

Employee

Name	Department	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

```
SELECT *
FROM Employee
WHERE Department IN
('Marketing',
'Sales');
```

43

IN

Employee

Name	Department	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

```
SELECT *
FROM Employee
WHERE Department IN
('Marketing',
'Sales');
```

Employee

Name	Department	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane

44

IN

Employee

Name	Department	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

```
SELECT *
FROM Employee
WHERE Department =
'Marketing' OR
Department = 'Sales';
```

Employee

Name	Department	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane

45

NOT IN

Employee

Name	Department	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

```
SELECT *
FROM Employee
WHERE Name NOT IN
(SELECT Manager
FROM Employee);
```

46

NOT IN

- First the subquery
SELECT Manager
FROM Employee
- is evaluated giving

Manager
Chris
Chris
Jane
Jane

- This gives
SELECT *
FROM Employee
WHERE Name NOT
IN ('Chris',
'Jane');

Name	Department	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Peter	Sales	Jane

47

NOT IN (depending on table def.)

Employee

Name	Department	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

```
SELECT *
FROM Employee
WHERE Name NOT IN
(SELECT Manager
FROM Employee);
```

Name	Department	Manager
------	------------	---------

48

Correct query

Employee

Name	Department	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

```
SELECT * FROM Employee
WHERE Name NOT IN
(SELECT Manager FROM
Employee WHERE
Manager IS NOT NULL);
```

49

EXISTS

- Using EXISTS we can see whether there is at least one element in a given set
- NOT EXISTS is true if the set is empty
- The set is always given by a subquery

```
SELECT columns
FROM tables
WHERE EXISTS set;
```

```
SELECT columns
FROM tables
WHERE NOT EXISTS
set;
```

50

EXISTS

Employee

Name	Dept	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

```
SELECT *
FROM Employee AS E1
WHERE EXISTS (
SELECT * FROM
Employee AS E2
WHERE E1.Name =
E2.Manager);
```

- Retrieve all the info for those employees who are also managers.

51

EXISTS

```
SELECT * FROM Employee AS E1 WHERE EXISTS
(SELECT * FROM Employee AS E2 WHERE E1.Name = E2.Manager);
```

Employee E1

Name	Dept	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

Employee E2

Name	Dept	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

52

EXISTS

Name	Dept	Manager	Name	Dept	Manager
John	Marketing	Chris	John	Marketing	Chris
John	Marketing	Chris	Mary	Marketing	Chris
John	Marketing	Chris	Chris	Marketing	Jane
John	Marketing	Chris	Peter	Sales	Jane
John	Marketing	Chris	Jane	Management	Jane
Mary	Marketing	Chris	John	Marketing	Chris
Mary	Marketing	Chris	Mary	Marketing	Chris
Mary	Marketing	Chris	Chris	Marketing	Jane
Mary	Marketing	Chris	Peter	Sales	Jane
Mary	Marketing	Chris	Jane	Management	Jane
Chris	Marketing	Jane	John	Marketing	Chris
Chris	Marketing	Jane	Mary	Marketing	Chris
Chris	Marketing	Jane	Chris	Marketing	Jane

53

EXISTS

```
SELECT * FROM Employee AS E1 WHERE EXISTS
(SELECT * FROM Employee AS E2 WHERE E1.Name = E2.Manager);
```

Name	Dept	Manager	Name	Dept	Manager
Chris	Marketing	Jane	John	Marketing	Chris
Chris	Marketing	Jane	Mary	Marketing	Chris
Jane	Management		Chris	Marketing	Jane
Jane	Management		Peter	Sales	Jane

Name	Dept	Manager
Chris	Marketing	Jane
Chris	Marketing	Jane
Jane	Management	
Jane	Management	

54

EXISTS

Employee

Name	Dept	Manager
John	Marketing	Chris
Mary	Marketing	Chris
Chris	Marketing	Jane
Peter	Sales	Jane
Jane	Management	

```
SELECT *
FROM Employee AS E1
WHERE EXISTS (
  SELECT * FROM
  Employee AS E2
  WHERE E1.Name =
  E2.Manager);
```

Name	Dept	Manager
Chris	Marketing	Jane
Jane	Management	

55

ANY and ALL

- ANY and ALL compare a single value to a set of values
- They are used with comparison operators like =, >, <, <>, >=, <=
- **val = ANY (set)** is true if there is at least one member of the set equal to value
- **val = ALL (set)** is true if all members of the set are equal to the value

56

ALL

Name	Salary
Mary	20,000
John	15,000
Jane	25,000
Paul	30,000

- Find the name(s) of the employee(s) who earn the highest salary

57

ALL

Name	Salary
Mary	20,000
John	15,000
Jane	25,000
Paul	30,000

- Find the name(s) of the employee(s) who earn the highest salary

Name
Paul

```
SELECT Name
FROM Employee
WHERE Salary >=
  ALL (
    SELECT Salary
    FROM Employee);
```

58

ANY

Name	Salary
Mary	20,000
John	15,000
Jane	25,000
Paul	30,000

- Find the name(s) of the employee(s) who earn more than someone else

59

ANY

Name	Salary
Mary	20,000
John	15,000
Jane	25,000
Paul	30,000

- Find the name(s) of the employee(s) who earn more than someone else

```
SELECT Name
FROM Employee
WHERE Salary >
  ANY (
    SELECT Salary
    FROM Employee);
```

60

Word Search

- Word Search
 - Commonly used for searching product catalogues etc.
 - Need to search by keywords
 - Might need to use partial keywords
- For example: Given a database of books, searching for "crypt" might return
 - "Cryptonomicon" by Neil Stephenson
 - "Applied Cryptographer" by Bruce Schneier

61

LIKE

- We can use the **LIKE** keyword to perform string comparisons in queries
- Like is not the same as '=' because it allows wildcard characters
- It is NOT normally case sensitive

```
SELECT * FROM books
WHERE bookName LIKE "%crypt%";
```

62

LIKE

- The '%' character can represent any number of characters, including none
- The '_' character represents exactly one character

```
bookName LIKE "crypt%"      bookName LIKE "cloud_"
```

- Will return "Cryptography Engineering" and "Cryptonomicon" but not "Applied Cryptography"
- Will return "Clouds" but not "Cloud" or "cloud computing"

63

LIKE

- Sometimes you might need to search for a set of words
 - To find entries with all words you can link conditions with AND
 - To find entries with any words use OR

```
SELECT * FROM
books
WHERE bookName
LIKE "%crypt%"
OR bookName LIKE
"%cloud%";
```

64

Example

Track				
cdID	Num	Track_title	Time	alID
1	1	Violent	239	1
1	2	Every Girl	410	1
1	3	Breather	217	1
1	4	Part of Me	279	1
2	1	Star	362	1
2	2	Teaboy	417	2

CD		
cdID	Title	Price
1	Mix	9.99
2	Compilation	12.99

Artist	
alID	Name
1	Stellar
2	Cloudboy

Write a query to find any track title containing either the string 'boy' or 'girl'

65

Example

```
SELECT Track_title FROM Track
WHERE Track_title LIKE "%boy%"
OR Track_title LIKE "%girl%";
```

66

Date, Time, Datetime

- 3 different types for a time column

Type	Description	Example
DATE	A Day, Month and Year	'1981-12-16' or '81-12-16'
TIME	Hours, Minutes and Seconds	'15:24:39'
DATETIME	Combination of above	'1981-12-16 15:24:39'

- Timestamp: as Datetime, usually used to display current date and time ('2014-11-04 15:30:43')
- Usual conditions may be used on WHERE clauses
 - `SELECT * FROM table-name`
`WHERE date-of-event < '2012-01-01' ;`
 - Or `WHERE date-of-event LIKE '2014-11-%' ;`

67

Take home messages

1. SELECT query to retrieve information
2. WHERE clause to specify condition
3. Cartesian product to combine tables
4. Different ways to write SELECT queries
 - a. Some more elegant than others

68

Thanks for your attention!

Any questions??

69